

```
import numpy as np
import pandas as pd
```

```
!iconv -f ISO-8859-1 -t UTF-8 spam.csv > spam_utf8.csv
```

```
df = pd.read_csv('spam_utf8.csv')
```

```
df.sample(5)
```

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
2876	ham	twenty past five he said will this train have ...	NaN	NaN	NaN
3807	ham	Mm you ask him to come its enough :-)	NaN	NaN	NaN
1769	ham	Ha... Both of us doing e same thing. But i got...	NaN	NaN	NaN
4012	ham	Ok.	NaN	NaN	NaN
373	ham	I cant keep talking to people if am not sure i...	NaN	NaN	NaN

```
df.shape
```

(5572, 5)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0    v1          5572 non-null   object
1    v2          5572 non-null   object
2    Unnamed: 2   50 non-null     object
3    Unnamed: 3   12 non-null     object
4    Unnamed: 4   6 non-null      object
dtypes: object(5)
memory usage: 217.8+ KB
```

```
df.drop(columns=['Unnamed: 2','Unnamed: 3','Unnamed: 4'],inplace=True)
```

```
df.sample(5)
```

1 to 5 of 5 entries Filter ?

index	v1	v2
4465	ham	Hey u still at the gym?
5252	ham	You do your studies alone without anyones help. If you cant no need to study.
3676	ham	Great! So what attracts you to the brothas?
3194	ham	Great. P diddy is my neighbor and comes for toothpaste every morning
890	ham	Why do you ask princess?

Show 25 per page



Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.

```
df.rename(columns={'v1': 'target', 'v2': 'text'},inplace=True)
df.sample(5)
```

	target	text
170	ham	Sir, I need AXIS BANK account no and bank addr...
989	ham	Ugh. Gotta drive back to sd from la. My butt i...
2534	ham	Ok enjoy . R u there in home.
225	ham	Would really appreciate if you call me. Just n...
756	ham	Cant think of anyone with * spare room off * t...

```
from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()
```

```
df['target'] = encoder.fit_transform(df['target'])
```

```
df.head()
```

	target	text
0	0	Go until jurong point, crazy.. Available only ...
1	0	Ok lar... Joking wif u oni...
2	1	Free entry in 2 a wkly comp to win FA Cup fina...
3	0	U dun say so early hor... U c already then say...
4	0	Nah I don't think he goes to usf, he lives aro...

```
df.isnull().sum()
```

```
target    0
text      0
dtype: int64
```

```
df.duplicated().sum()
```

```
403
```

```
df = df.drop_duplicates(keep='first')
```

```
df.duplicated().sum()
```

```
0
```

```
df.shape
```

```
(5169, 2)
```


```
df.head()
```

	target	text
0	0	Go until jurong point, crazy.. Available only ...
1	0	Ok lar... Joking wif u oni...
2	1	Free entry in 2 a wkly comp to win FA Cup fina...
3	0	U dun say so early hor... U c already then say...
4	0	Nah I don't think he goes to usf, he lives aro...

```
df['target'].value_counts()
```

```
0    4516
1     653
Name: target, dtype: int64
```

```
import matplotlib.pyplot as plt
plt.pie(df['target'].value_counts(), labels=['ham', 'spam'], autopct="%0.2f")
plt.show()
```



```
import nltk
```

```
!pip install nltk
```

Requirement already satisfied: nltk in /usr/local/lib/python3.10/dist-packages (3.8.1)
Requirement already satisfied: click in /usr/local/lib/python3.10/dist-packages (from nltk) (8.1.7)
Requirement already satisfied: joblib in /usr/local/lib/python3.10/dist-packages (from nltk) (1.3.2)
Requirement already satisfied: regex<=2021.8.3 in /usr/local/lib/python3.10/dist-packages (from nltk) (2023.6.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from nltk) (4.66.1)

```
nltk.download('punkt')
```

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip.
True

```
df['num_characters'] = df['text'].apply(len)
```

```
df.head()
```

target		text	num_characters	
0	0	Go until jurong point, crazy.. Available only ...	111	
1	0	Ok lar... Joking wif u oni...	29	
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	
3	0	U dun say so early hor... U c already then say...	49	
4	0	Nah I don't think he goes to usf, he lives aro...	61	



```
df['num_words'] = df['text'].apply(lambda x:len(nltk.word_tokenize(x)))
```

```
df.head()
```



target		text	num_characters	num_words	
0	0	Go until jurong point, crazy.. Available only ...	111	24	
1	0	Ok lar... Joking wif u oni...	29	8	
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	37	
3	0	U dun say so early hor... U c already then say...	49	13	
4	0	Nah I don't think he goes to usf, he lives aro...	61	15	

```
df['num_sentences'] = df['text'].apply(lambda x:len(nltk.sent_tokenize(x)))
```



```
df.head()
```

	target	text	num_characters	num_words	num_sentences	
0	0	Go until jurong point, crazy.. Available only ...	111	24	2	
1	0	Ok lar... Joking wif u oni...	29	8	2	
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	37	2	
3	0	U dun say so early hor... U c already then sav...	49	13	1	



```
df[['num_characters', 'num_words', 'num_sentences']].describe()
```

	num_characters	num_words	num_sentences	
count	5169.000000	5169.000000	5169.000000	
mean	78.977945	18.455794	1.965564	
std	58.236293	13.324758	1.448541	
min	2.000000	1.000000	1.000000	
25%	36.000000	9.000000	1.000000	
50%	60.000000	15.000000	1.000000	
75%	117.000000	26.000000	2.000000	
max	910.000000	220.000000	38.000000	

```
df[df['target'] == 0][['num_characters', 'num_words', 'num_sentences']].describe()
```

	num_characters	num_words	num_sentences	
count	4516.000000	4516.000000	4516.000000	
mean	70.459256	17.123782	1.820195	
std	56.358207	13.493970	1.383657	
min	2.000000	1.000000	1.000000	
25%	34.000000	8.000000	1.000000	
50%	52.000000	13.000000	1.000000	
75%	90.000000	22.000000	2.000000	
max	910.000000	220.000000	38.000000	


```
df[df['target'] == 1][['num_characters', 'num_words', 'num_sentences']].describe()
```

	num_characters	num_words	num_sentences	
count	653.000000	653.000000	653.000000	
mean	137.891271	27.667688	2.970904	
std	30.137753	7.008418	1.488425	
min	13.000000	2.000000	1.000000	
25%	132.000000	25.000000	2.000000	
50%	149.000000	29.000000	3.000000	
75%	157.000000	32.000000	4.000000	
max	224.000000	46.000000	9.000000	

```
import seaborn as sns
```


```
plt.figure(figsize=(12,6))
sns.histplot(df[df['target'] == 0]['num_characters'])
sns.histplot(df[df['target'] == 1]['num_characters'],color='red')
```

<Axes: xlabel='num_characters', ylabel='Count'>




```
plt.figure(figsize=(12,6))
sns.histplot(df[df['target'] == 0]['num_words'])
sns.histplot(df[df['target'] == 1]['num_words'],color='red')
```

<Axes: xlabel='num_words', ylabel='Count'>

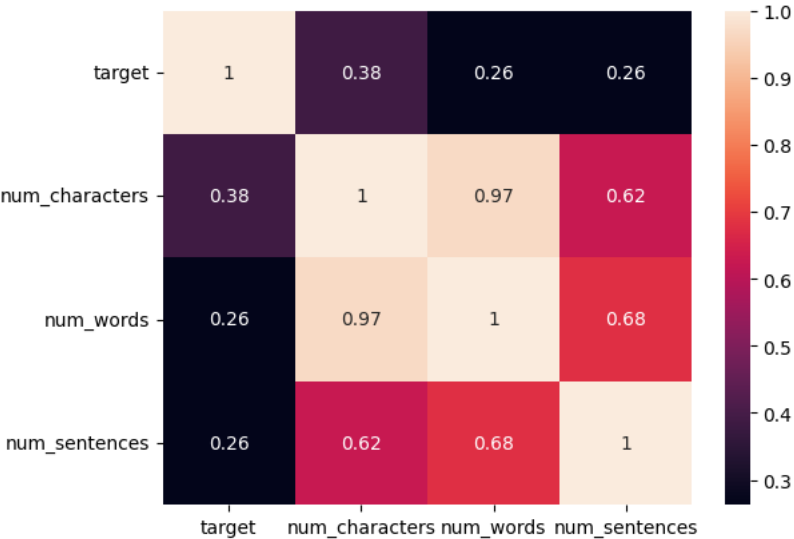


```
sns.pairplot(df,hue='target')
```



```
sns.heatmap(df.corr(),annot=True)
```

<ipython-input-38-8df7bcac526d>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future ver
sns.heatmap(df.corr(),annot=True)
<Axes: >



```
def transform_text(text):
    text = text.lower()
    text = nltk.word_tokenize(text)

    y = []
    for i in text:
        if i.isalnum():
            y.append(i)

    text = y[:]
    y.clear()

    for i in text:
        if i not in stopwords.words('english') and i not in string.punctuation:
            y.append(i)

    text = y[:]
    y.clear()

    for i in text:
        y.append(ps.stem(i))

    return " ".join(y)
```

```
!pip install nltk
```

```
Requirement already satisfied: nltk in /usr/local/lib/python3.10/dist-packages (3.8.1)
Requirement already satisfied: click in /usr/local/lib/python3.10/dist-packages (from nltk) (8.1.7)
Requirement already satisfied: joblib in /usr/local/lib/python3.10/dist-packages (from nltk) (1.3.2)
Requirement already satisfied: regex<=2021.8.3 in /usr/local/lib/python3.10/dist-packages (from nltk) (2023.6.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from nltk) (4.66.1)
```

```
!nltk.download('stopwords')
```

```
/bin/bash: -c: line 1: syntax error near unexpected token `stopwords'
/bin/bash: -c: line 1: `nltk.download('stopwords')'
```

```
from nltk.corpus import stopwords
stopwords.words('english')
```

```
'needn',
"needn't",
'shan',
"shan't",
'shouldn',
"shouldn't",
'wasn',
"wasn't",
'weren',
"weren't",
'won',
"won't",
'wouldn',
"wouldn't"]
```

```
import string
string.punctuation
```

```
'!"#$%&\'()*+,-./:;<=>@[\\]^_`{|}~'
```

```
df['text'][10]
```

```
'I'm gonna be home soon and i don't want to talk about this stuff anymore tonight, k? I've cried enough today.'
```

```
from nltk.stem.porter import PorterStemmer
ps = PorterStemmer()
ps.stem('loving')
```

```
'love'
```

```
df['transformed_text'] = df['text'].apply(transform_text)
```

```
df.head()
```

	target	text	num_characters	num_words	num_sentences	transformed_text
0	0	Go until jurong point, crazy.. Available only ...	111	24	2	go jurong point crazi avail bugi n great world...
1	0	Ok lar... Joking wif u oni...	29	8	2	ok lar joke wif u oni
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	37	2	free entri 2 wkli comp win fa cup final tkt 21...
3	0	U dun say so early hor... U c already then sav...	49	13	1	u dun say earli hor u c already say

```
from wordcloud import WordCloud
wc = WordCloud(width=500,height=500,min_font_size=10,background_color='white')
```

```
spam_wc = wc.generate(df[df['target'] == 1]['transformed_text'].str.cat(sep=" "))
```

```
plt.figure(figsize=(15,6))
plt.imshow(spam_wc)
```



```
plt.figure(figsize=(15,6))
plt.imshow(ham_wc)
```

[illegible]

<https://colab.research.google.com/drive/1V19wnQYBrqy6UdirYPii4BtIXSnZZMK3?authuser=0#scrollTo=Y8nNQXLehhpB&uniqifier=2&printMod...> 9/14


index	target	text	num_characters	num_words	num_sentences	transformed_text
0	0	Go until jurong point, crazy.. Available only in bugis n great world la e buffet... Cine there got amore wat...	111	24	2	go jurong point crazi avail bugi n great world la e buffet cine got amor wat
1	0	Ok lar... Joking wif u oni...	29	8	2	ok lar joke wif u oni
2	1	Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005. Text FA to 87121 to receive entry question(std txt rate)T&C's apply 08452810075over18's	155	37	2	free entri 2 wkli comp win fa cup final tkt 21st may text fa 87121 receiv entri question std txt rate c appli 08452810075over18
3	0	U dun say so early hor... U c already then say...	49	13	1	u dun say earli hor u c already say
4	0	Nah I don't think he goes to usf, he lives around here though	61	15	1	nah think goe usf live around though

Show 25 per page




Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.


Distributions




Categorical distributions




2-d distributions




Time series



Values



2-d categorical distributions



Faceted distributions

```
spam_corpus = []
for msg in df[df['target'] == 1]['transformed_text'].tolist():
    for word in msg.split():
        spam_corpus.append(word)
```

```
len(spam_corpus)
```

9939

```
# Text Vectorization
# using Bag of Words
df.head()
```

	target	text	num_characters	num_words	num_sentences	transformed_text
0	0	Go until jurong point, crazy.. Available only ...	111	24	2	go jurong point crazi avail bugi n great world...
1	0	Ok lar... Joking wif u oni...	29	8	2	ok lar joke wif u oni
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	37	2	free entri 2 wkli comp win fa cup final tkt 21...
3	0	U dun say so early hor... U c already then sav...	49	13	1	u dun say earli hor u c already say

```
from sklearn.feature_extraction.text import CountVectorizer,TfidfVectorizer
cv = CountVectorizer()
tfidf = TfidfVectorizer(max_features=3000)
```

```
X = tfidf.fit_transform(df['transformed_text']).toarray()
```

```
X.shape

(5169, 3000)
```

```
y = df['target'].values
```

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=2)
```

```
from sklearn.naive_bayes import GaussianNB,MultinomialNB,BernoulliNB
from sklearn.metrics import accuracy_score,confusion_matrix,precision_score
```

```
gnb = GaussianNB()
mnb = MultinomialNB()
bnb = BernoulliNB()
```

```
gnb.fit(X_train,y_train)
y_pred1 = gnb.predict(X_test)
print(accuracy_score(y_test,y_pred1))
print(confusion_matrix(y_test,y_pred1))
print(precision_score(y_test,y_pred1))
```

```
0.8694390715667312
[[788 108]
 [ 27 111]]
0.5068493150684932
```

```
mnb.fit(X_train,y_train)
y_pred2 = mnb.predict(X_test)
print(accuracy_score(y_test,y_pred2))
print(confusion_matrix(y_test,y_pred2))
print(precision_score(y_test,y_pred2))
```

```
0.9709864603481625
[[896  0]
 [ 30 108]]
1.0
```

```
bnb.fit(X_train,y_train)
y_pred3 = bnb.predict(X_test)
print(accuracy_score(y_test,y_pred3))
print(confusion_matrix(y_test,y_pred3))
print(precision_score(y_test,y_pred3))
```

```
0.9835589941972921
[[895  1]
 [ 16 122]]
0.991869918699187
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier

```

```

svc = SVC(kernel='sigmoid', gamma=1.0)
knc = KNeighborsClassifier()
mnb = MultinomialNB()
dtc = DecisionTreeClassifier(max_depth=5)
lrc = LogisticRegression(solver='liblinear', penalty='l1')
rfc = RandomForestClassifier(n_estimators=50, random_state=2)
abc = AdaBoostClassifier(n_estimators=50, random_state=2)
bc = BaggingClassifier(n_estimators=50, random_state=2)
etc = ExtraTreesClassifier(n_estimators=50, random_state=2)
gbdt = GradientBoostingClassifier(n_estimators=50, random_state=2)
xgb = XGBClassifier(n_estimators=50, random_state=2)

```

```

clfs = {
    'SVC' : svc,
    'KN' : knc,
    'NB': mnb,
    'DT': dtc,
    'LR': lrc,
    'RF': rfc,
    'AdaBoost': abc,
    'BgC': bc,
    'ETC': etc,
    'GBDT': gbdt,
    'xgb': xgb
}

```

```

def train_classifier(clf,X_train,y_train,X_test,y_test):
    clf.fit(X_train,y_train)
    y_pred = clf.predict(X_test)
    accuracy = accuracy_score(y_test,y_pred)
    precision = precision_score(y_test,y_pred)

    return accuracy,precision

```

```

train_classifier(svc,X_train,y_train,X_test,y_test)

(0.9758220502901354, 0.9747899159663865)

```

```

accuracy_scores = []
precision_scores = []

for name,clf in clfs.items():

    current_accuracy,current_precision = train_classifier(clf, X_train,y_train,X_test,y_test)

    print("For ",name)
    print("Accuracy - ",current_accuracy)
    print("Precision - ",current_precision)

    accuracy_scores.append(current_accuracy)
    precision_scores.append(current_precision)

```

```




For SVC
Accuracy - 0.9758220502901354
Precision - 0.9747899159663865
For KN
Accuracy - 0.9052224371373307
Precision - 1.0
For NB
Accuracy - 0.9709864603481625
Precision - 1.0
For DT
Accuracy - 0.9303675048355899
Precision - 0.8173076923076923
For LR
Accuracy - 0.9584139264990329
Precision - 0.9702970297029703
For RF

```

```
Accuracy - 0.9758220502901354
Precision - 0.9829059829059829
For AdaBoost
Accuracy - 0.960348162475822
Precision - 0.9292035398230089
For BgC
Accuracy - 0.9584139264990329
Precision - 0.8682170542635659
For ETC
Accuracy - 0.9748549323017408
Precision - 0.9745762711864406
For GBDT
Accuracy - 0.9468085106382979
Precision - 0.9191919191919192
For xgb
Accuracy - 0.9671179883945842
Precision - 0.9262295081967213
```

```
performance_df = pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy':accuracy_scores,'Precision':precision_scores}).sort_values('Precision')
```

performance_df

	Algorithm	Accuracy	Precision	
1	KN	0.905222	1.000000	  
2	NB	0.970986	1.000000	
5	RF	0.975822	0.982906	
0	SVC	0.975822	0.974790	
8	ETC	0.974855	0.974576	
4	LR	0.958414	0.970297	
6	AdaBoost	0.960348	0.929204	
10	xgb	0.967118	0.926230	
9	GBDT	0.946809	0.919192	
7	BgC	0.958414	0.868217	
3	DT	0.930368	0.817308	

```
performance_df1 = pd.melt(performance_df, id_vars = "Algorithm")
```

performance_df1

	Algorithm	variable	value
0	KN	Accuracy	0.905222
4	NP	Accuracy	0.970088

```
sns.catplot(x = 'Algorithm', y='value',  
            hue = 'variable',data=performance_df1, kind='bar',height=5)  
plt.ylim(0.5,1.0)  
plt.xticks(rotation='vertical')  
plt.show()
```

